

Assignment 2

TCP Socket Programming

Prof. Ai-Chun Pang

TA / Yu-Yu Chen, Shao-Cheng Fan, Kuang-Hui Huang

Assignment 2 Announcement

Specification (1/22)

- In this assignment, you need to implement a simple Network Storage System with the following functions:
 - client can **list all files in the user folder of the server**.
 - client can **upload files** to the server.
 - client can **download files** from the server.
 - client can **watch a “.mpg” video (streaming)** on the server.
 - User “admin” can **ban users**.
 - User “admin” can **unban users**.
 - User “admin” can **list all banned users**.

Specification (2/22)

- For video streaming, you **don't need to send audio**. You can just send frames in **RAW format**.
- After upload and download, you must **ensure the files are identical between source and destination**.
- In this assignment, all the transmissions should be implemented in **C/C++** and by **TCP socket**.

Specification – Commands (3/22)

- The server is required to support multiple connections. That is, **more than one client** can connect to the server **simultaneously**.
- After building up a connection, a **client** can use the commands below:

```
$ ls
```

```
$ put <filename>
```

```
$ get <filename>
```

```
$ play <videofile.mpg>
```

Specification – Admin Commands (4/22)

- After building up a connection, user “admin” can use **the basic commands and the commands below**:

```
$ ban <username1> <username2> ... <usernameN>
```

```
$ unban <username1> <username2> ... <usernameN>
```

```
$ blocklist
```

- Don't assume the lengths of the commands are lower than a specific size in advance.

Specification – ls (5/22)

- For the command **ls**, a client should list all the files in the user folder of the server with the filename separated by a **newline** character.

```
$ ls  
file1  
file2  
file3  
...  
video.mpg
```

Specification – put (6/22)

- For the command **put**, a client should upload the file specified in the arguments to the server.

```
$ put file1
```

```
putting file1...
```

```
$ put file99
```

```
file99 doesn't exist.
```

- There won't be multiple clients putting the same file at the same time.

Specification – get (7/22)

- For the command **get**, a client should download the file specified in the arguments from the server.

```
$ get file1
```

```
getting file1...
```

```
$ get file99
```

```
file99 doesn't exist.
```

- There might be multiple clients getting the same file at the same time.
- A client will not get the file that is being put by another client.

Specification – play (8/22)

- For the command **play**, a client should play the **.mpg** video file specified in the arguments.

```
$ play video1.mpg  
playing the video...  
$ play video99.mpg  
video99.mpg doesn't exist.
```

- There might be multiple clients playing the same video at the same time.
- A client will not play the video that is being put.

Specification – play (9/22)

- After pressing **ESC**, a client should terminate the video and be able to keep sending another command.
- The video window should be closed automatically at the end of the video and after pressing ESC.

Specification – ban (10/22)

- User “admin” can add other users to the **blocklist** by the **ban** command.

```
$ ban Alice Bob Cat admin Bob
```

```
Ban Alice successfully!
```

```
Ban Bob successfully!
```

```
Ban Cat successfully!
```

```
You cannot ban yourself!
```

```
User Bob is already on the blocklist!
```

```
$ ban Cat
```

```
User Cat is already on the blocklist!
```

Specification – ban (cont.) (11/22)

- User “admin” can ban many users in one command.
- There may be no connection between the banned users and the server.
That is, you don't need to check whether those banned users connect to the server.

Specification – unban (12/22)

- User “admin” can remove other users from the **blocklist** by the **unban** command.

```
$ unban Alice Peter admin Alice
```

```
Successfully removed Alice from the blocklist!
```

```
User Peter is not on the blocklist!
```

```
User admin is not on the blocklist!
```

```
User Alice is not on the blocklist!
```

- User “admin” can unban many users in one command.

Specification – blocklist (13/22)

- For the command **blocklist**, user “admin should list every banned user with the username separated by a **newline** character.

```
$ blocklist
```

```
Bob
```

```
Cat
```

- If the blocklist is empty, you don't need to print anything.

Specification – Permission (14/22)

- Only user “admin” can run **ban**, **unban**, and **blocklist**. If other users run the commands, print out “**Permission denied.**” on the client side.
- The users on the blocklist can't run any commands, and they should print out “**Permission denied.**”
- If users are banned while running a command, they can keep running until the end of the command.

Specification – Permission (cont.) (15/22)

- User “Bob” was added to the blocklist when playing the video.
- The sample output of user “Bob”:

```
$ blocklist
```

```
Permission denied.
```

```
$ play video1.mpg
```

```
playing the video...
```

```
// admin added user Bob to the blocklist.
```

```
$ play video1.mpg
```

```
Permission denied.
```

Specification – Compilation (16/22)

- You are required to write a Makefile for compilation.

```
$ make client      // To compile client code  
$ make server      // To compile server code
```

- After compilation, there will be 2 binary files named “client” and “server.”

Specification – Usage (17/22)

- When we launch the server, we will use

```
$ ./server <port>           // <port> will be determined
```

- After launching, the server should create a folder named **server_dir**.
- No matter what happens, **server** should NOT be terminated.

Specification – Usage (18/22)

- When we launch the client, we will use

```
$ ./client <username> <ip>:<port>
```

```
// <username> consists of no more than 10 English letters and digits
```

```
// <ip> is the ip address of the server
```

```
// <port> is determined by the command above
```

- After launching, the client should create a folder named **client_dir**, and the server should create a folder named **<username>** under folder **server_dir**.
- A client might be terminated at any time.

Specification – Usage (19/22)

- If we launch the server and 2 clients with username **Alice** and **Bob**:

```
|-- server_dir
|  |-- Alice
|  |  |-- file00
|  |  `-- file01
|  `-- Bob
|      |-- file10
|      `-- file11
`-- client_dir
```

Specification – Error Handling (20/22)

- All of your outputs should be printed to **standard output (stdout)** and end with a **newline**.
- If the command doesn't exist or the command format is wrong, print out "**Command not found.**" on the client side.
- If the file doesn't exist while putting, getting, or playing a file, print out "**<filename> doesn't exist.**" on the client side.
- If the video file is not a ".mpg" file while playing a video file, print out "**<videofile> is not an mpg file.**" on the client side.

Specification (21/22)

- The server should **output the file descriptor of the new connection** and **the username** when a client connects to the server.

Accept a new connection on socket [<file descriptor>]. Login as <username>.

- A client should be able to send another command after a command is finished.
- The multiple connections should be implemented with **pthread.h** or **select()**, while the video player should be implemented with **OpenCV**.
- The implementation must be in **C or C++**.
- The file size will be no more than 2GB.
- Every command should be done in 5 minutes.

Specification – Multiple Connections (22/22)

- There **can be more than one client** connecting to the server and sending commands to the server **simultaneously**.
- A single user can **log in using multiple clients simultaneously**.
- Commands from different **clients** should be executed **concurrently**.

That is, a command from one client **cannot** be blocked by other commands from other clients.

Grading Policy (1/4)

- This assignment accounts for 15% of the total score.
- **Command Sending (10%)**
 - The client and the server handle commands correctly. (5%)
 - The client prints out responses correctly. (5%)
- **Basic File Transferring (20%)**
 - There would be 5 test cases. (4% * 5)
 - You will get **0 points** in a test case if the transfer of a test case **is terminated** or **halts** before finished, or the files are **not identical between source and destination** after the transfer.

Grading Policy (2/4)

- **User Management (12%)**

- There will be 4 test cases. (3% * 4)

- **Video Streaming (18%)**

- Correctly playing a **resolution-fixed** video. (960*540) (6%)
- Correctly playing a **resolution-unknown** video. (6%)
(The client has no idea about the resolution of the video.)
- Playing video while others are transmitting files. (6%)

Grading Policy (3/4)

- **Multiple Connections (30%)**

- There will be 5 test cases.
- You need to implement one of the methods below and get the corresponding score for this part.
 - **Method 1:** Use `<pthread.h>` to achieve this function (basic) (20%)
 - **Method 2:** Use `select()` to achieve this function (advance) (30%)

Grading Policy (4/4)

- **Report** **(5% * 4)**
 - Draw a flowchart of the file transferring and explain how it works in detail.
 - Draw a flowchart of the video streaming and explain how it works in detail.
 - What is **SIGPIPE**? Is it possible that **SIGPIPE** is sent to your process? If so, how do you handle it?
 - Is blocking I/O equal to synchronized I/O? Please give some examples to explain it.

Submission

- Requirements
 - Your report can be a **pdf** or **clear image** file. Submit it to **Gradescope**.
 - Please put all the source code (i.e., without your report, the video file, and the execution file) into a folder named **<studentID>_hw2** and compress the folder as a **.zip** file. Submit your **.zip** file to **NTU COOL**.

```
B09902999_hw2.zip
```

```
`-- B09902999_hw2          (== folder)
```

```
    |-- your source code and Makefile
```

- The penalty for the wrong format is **10 points**.
- **No plagiarism is allowed. A plagiarist will be graded zero.**

Submission

- Deadline
 - Due Date : 23:59:00, November 8th, 2022
 - The penalty for late submission is **20 points per day**.

If You have any Problems...

- You can
 - Ask questions on NTU COOL Discussion Forum, or
 - Send a mail to TA with the tag **[HW2]** in the title
 - Ask questions in TA hours in R438 by appointment. [Google Sheet Link](#).
- TA Email: ntu.cnta@gmail.com

Sample Codes

- We will provide sample codes for your reference.
 - *server.c* - Default port number is 8787
 - *client.c* - Default IP address is 127.0.0.1, port is 8787
 - *pthread.c*
 - *openCV.cpp* - It will play video.mpg
 - *video.mpg*
 - *Makefile*

Environment Setup

Environment

- We provide a VirtualBox VM for you to run our example code. If you use computers with ARM-based processors (e.g., M1 Macbook), you can use NTU CSIE workstations alternatively.
- **Please make sure you can compile and run your code well on either our VM or NTU CSIE workstation.**
- If you use NTU CSIE workstations, please enable X11 forwarding in order to play videos.
 - For mac users, you can use XQuartz.
 - For Windows users, you can use MobaXterm.
 - Linux supports this feature natively.

Environment(cont.)

- If you choose to set up the environment on your OS rather than using our VM, here is information about our environment
 - Ubuntu 20.04 x64
 - OpenCV 4.2.0
- You can install OpenCV 4.2.0 by the command below or follow the instruction [here](#):

```
$ sudo apt install libopencv-dev=4.2.0+dfsg-5
```

VirtualBox Setup

- Download the VM from
 - Our Google Drive
 - The password of our VM is **cn2022**.
- Install Virtualbox (natively installed on the computers of Lab R204).

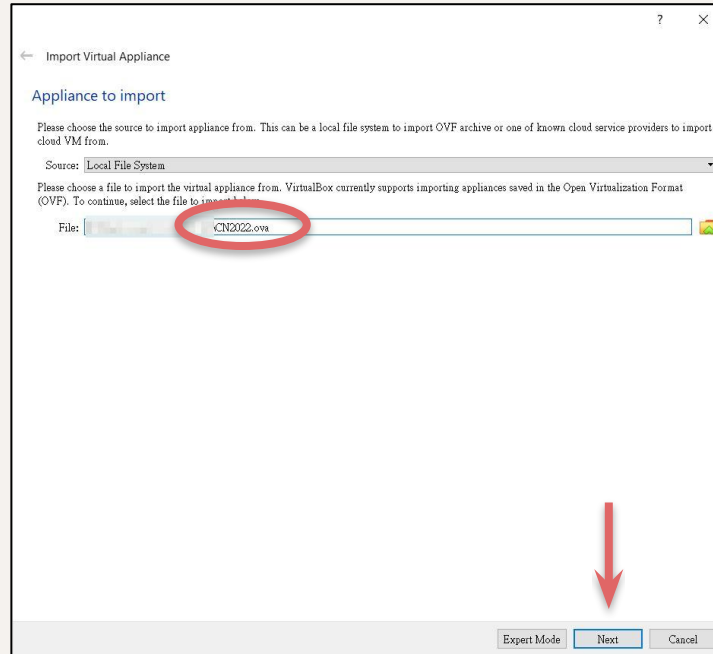
VirtualBox Setup

- Click “**Import**” to import the “**CN2022.ova**”



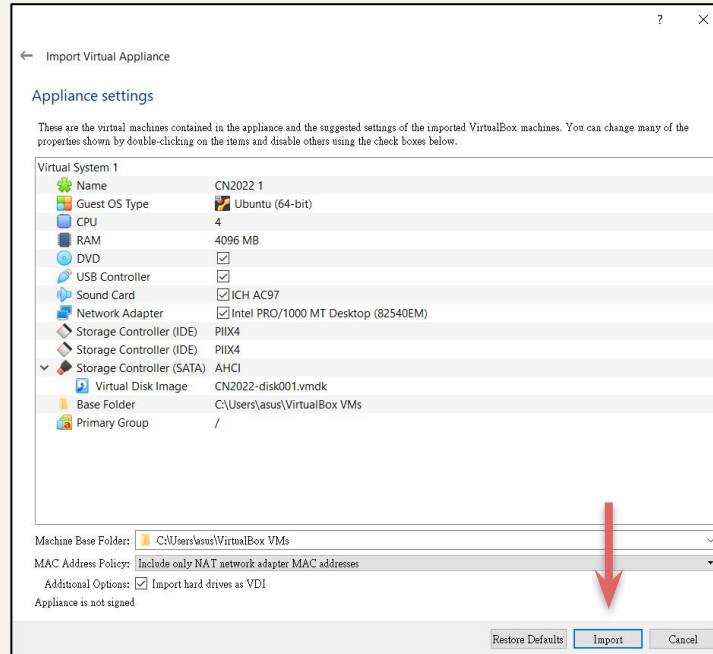
VirtualBox Setup

- Choose “**CN2022.ova**” file and click “**Next**”



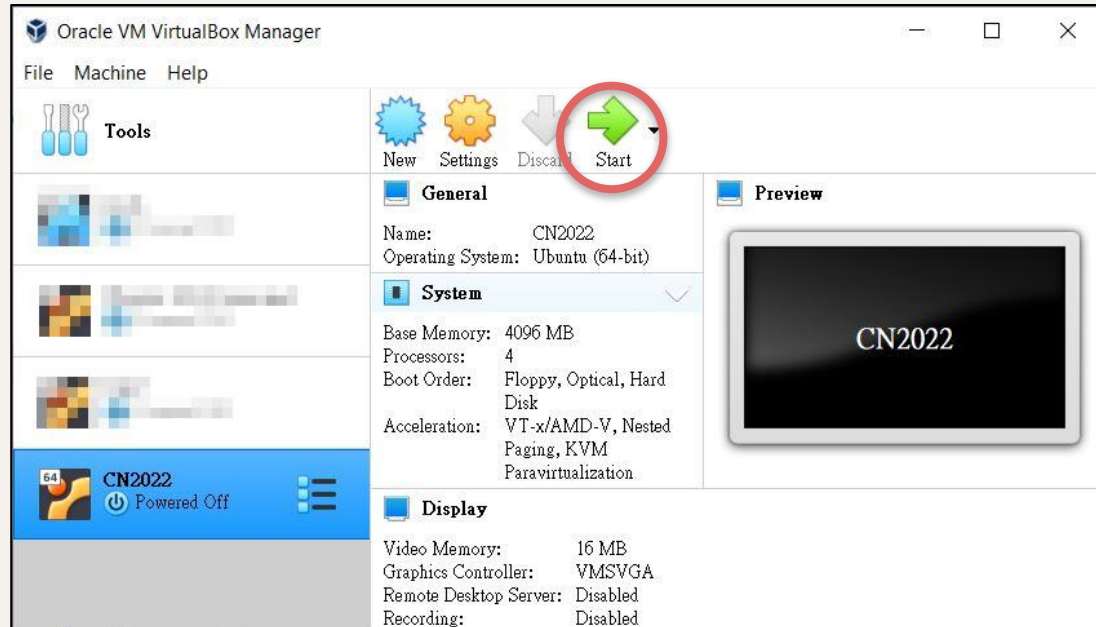
VirtualBox Setup

- Click “**Import**” to import the “**CN2022.ova**”



VirtualBox Setup

- Choose “**CN2022**” and then start the machine.



Auxiliary Libraries

OpenCV

- An open source library for computer vision.
- **Mat** is an image container to load an image so that you can easy to do image processing, recognition, etc.
- In this assignment, we use this library to get frames from videos on the server, and show frames on the client.
- We will provide a sample code and a .mpg file for you.
- To compile code with OpenCV,

```
$ g++ <file name> -o <output name> `pkg-config --cflags --libs opencv4`
```

Pthread


- Pthread, i.e., POSIX Thread, is used to implement multi-thread parallelization in a POSIX environment.
- You can use pthread to achieve multiple connections.
- You don't need to deal with synchronization issues, i.e., in our test cases, it won't put a file with the same file name.
- We will provide a sample code for you.
- To compile with Pthread,

```
$ g++ <file name> -o <output name> -pthread
```


Socket Programming Tutorial

What is Socket?

- **Socket** is the **API for the TCP/IP protocol** stack.
- Provides communication **between the Application layer and Transport layer**.
- Make internet communication **like a file descriptor**.
 - `read()` and `write()`
- We will provide a sample code for you.



What is Socket?



File Descriptors

- When we open an existing file or create a new file, the kernel returns a file descriptor to the process.
- If we want to read or write a file, we identify the file with the file descriptor.

Integer value	Name	<unistd.h> symbolic constant	<stdio.h> file stream
0	Standard input	STDIN_FILENO	stdin
1	Standard output	STDOUT_FILENO	stdout
2	Standard error	STDERR_FILENO	stderr

```
FILE *fp = fopen("this.txt","w");  
fprintf(fp, "Happy Coding.");  
fclose(fp);
```


```
printf("This is Computer Networking.\n");  
fprintf(stdout, " This is Computer Networking.\n");
```

TCP Service

- **TCP (Transmission Control Protocol)**
 - **Connection-oriented**
 - **Reliable transport**
 - Flow control
 - Congestion control
- What is Socket-Address?
 - **IP address + Port number**
 - IP address: To find out the machine (Network Layer)
 - Port number: To find out the process (Transport Layer)

TCP Socket Flow Chart

- Socket uses **blocking I/O** by default.



socket()

- **Create** the endpoint for connection.

```
#include <sys/socket.h>

int socket (int domain, int type, int protocol);
```

- **domain**
 - **AF_UNIX/AF_LOCAL**: communication between 2 processes on a host. So they can share a file system.
 - **AF_INET, AF_INET6**: communication between processes on different hosts through the Internet. **AF_INET** is for **IPv4**, whereas **AF_INET6** is for **IPv6**.

socket()

- **Create** the endpoint for connection.

```
#include <sys/socket.h>

int socket (int domain, int type, int protocol);
```

- **type**
 - **SOCK_STREAM**: sequential and connection-oriented (TCP)
 - **SOCK_DGRAM**: datagram (UDP)
- **protocol**: defined in /etc/protocols, usually set to 0
- **return**: socket file descriptor (an integer)

bind()

- **Bind** the address to the socket.

```
#include <sys/socket.h>

int bind (int sockfd, struct sockaddr *addr, socklen_t len);
```

- **sockfd**: specifies the socket file descriptor to bind.
- **addr**
 - specifies the socket address to be associated with the sockfd
 - You can use “**struct sockaddr_in***” defined in **<netinet/in.h>**, and then cast it into “**struct sockaddr***”
- **len**: specifies the size of sockaddr (= sizeof(struct sockaddr))

bind()

```
struct sockaddr_in {  
    short sin_family;           // address family. EX:AF_INET  
    unsigned short sin_port;    // port number for network  
    struct in_addr sin_addr;    // IP address for network  
    unsigned char sin_zero[8];  // pad to sizeof(struct sockaddr)  
}
```

listen()

- Specify a socket to **listen** for connections.

```
#include <sys/socket.h>

int listen (int sockfd, int backlog);    // returns 0 if it's success; -1 otherwise
```

- **sockfd**: specifies the socket file descriptor to listen.
- **backlog**: specifies the number of users allowed in queue.

accept()

- **Accept** the connection on a socket.

```
#include <sys/socket.h>

int accept (int sockfd, struct sockaddr *addr, socklen_t *addrlen);
```

- **sockfd**: specifies the socket being listened to.
- **addr**: pointer to the sockaddr. It will be filled in with **the address of the peer socket**.
- **Blocking** until a user connect() call is received.
- After accepting the connection, it **creates a new file descriptor** for the client. The original socket is not affected.

connect()

- **Connect** to the socket from client to server.

```
#include <sys/socket.h>
```

```
int connect (int sockfd, struct sockaddr *addr, socklen_t len);
```

- The format is the same as bind().

close()

- **Close the file descriptor.**

```
#include <unistd.h>  
int close (int sockfd);    // returns 0 if it's success; -1 otherwise
```

read()/recv()

- **Read** data from the socket file descriptor.

```
#include <unistd.h>

ssize_t read (int fd, void *buf, size_t count);
ssize_t recv (int fd, void *buf, size_t len, int flag);
```

- **fd**: specifies the socket file descriptor to read data from.
- **buf**: specifies the buffer to contain the received data.
- **count/len**: specifies the size to receive.
- **flag**: (**read()** has no this parameter.) It's about some details like blocking/nonblocking.

read()/recv()

- **Read** data from the socket file descriptor.

```
#include <unistd.h>

ssize_t read (int fd, void *buf, size_t count);
ssize_t recv (int fd, void *buf, size_t len, int flag);
```

- Reading data from a file may be
 - Successful, return the number of bytes received
 - EOF. (end of file) (i.e., return = 0)
 - Failed, **errno** is set to indicate the error.
- It may be blocked. (**block I/O**).

write()/send()

- **Write** data to socket file descriptor.

```
#include <unistd.h>

ssize_t write (int fd, const void *buf, size_t count);
ssize_t send (int fd, const void *buf, size_t len, int flags);
```

- **fd**: specifies the socket file descriptor to send data to.
- **buf**: specifies the buffer to contain the data to be transmitted.
- **count/len**: specifies the size to send.
- **flag**: (write()) has no this parameter.) It's about some details.

write()/send()

- **Write** data to socket file descriptor.

```
#include <unistd.h>

ssize_t write (int fd, const void *buf, size_t count);
ssize_t send (int fd, const void *buf, size_t len, int flags);
```

- Writing data to file may be
 - Successful, return the number of bytes written.
 - Failed, **errno** is set to indicate the error.
- It may be blocked. (**block I/O**)

Useful Functions

- Address and port numbers are stored as integers.
 - Different machines implement **different endian**.
 - They may communicate with each other on the network.
- Converting IP address and port number.
 - `htonl()`: for IP address (host -> network)
 - `ntohl()`: for IP address (network -> host)
 - `htons()`: for port number (host -> network)
 - `ntohs()`: for port number (network -> host)

Useful Functions

- An IP address is usually hard to remember.
 - We need to **translate the hostname to IP address**.
- Translate a hostname to IP address.

```
#include <netdb.h>

struct hostent *gethostbyname (const char *name);
```

Supplementary Materials

How to Emulate a Bad Network

- Some bugs may occur when the network is not good.
- We can emulate a bad network to test our program in advance.
- Linux
 - Get the network interfaces list in your machine

```
$ ifconfig
```

- Then, to make your personal internet slow, you can enter

```
$ sudo tc qdisc add dev <interface> root netem delay 500ms
```

In this way, the delay will be 500ms.

```
$ sudo tc qdisc del dev <interface> root netem
```

How to Trace Kernel

- Sometimes, the service may terminate without any error message.
- It happens usually because of some kernel issues.
- To trace the interactions between your code and kernel, we can use **strace**.
- To run your program with strace, you can enter

```
$ strace ./<name>
```

Behavior of `send()` and `recv()`

- In fact, a `send()` doesn't imply all the data in the buffer are sent **once**.
- In addition, a `send()` doesn't imply all the data in the buffer are sent in a packet, and even 2 `send()` don't imply they are in different packets.
- You are required to design a protocol so that each receive has the same size as the `send` in respect to it.

select()

- `select()` provides you to supervise multiple sockets, telling you which is able to read or write, etc.
- With `select()`, it is possible to achieve Asynchronous Blocking I/O.
- If you want to implement this assignment with `select()`, please refer to [this website](#).

select()

- **Monitor** whether there is at least one fd available.

```
#include <unistd.h>

int select (int nfds, fd_set*, readfds, fd_set* writefds, fd_set* exceptfds, struct
timeval* timeout);
```

- **nfds**: specifies the number of file descriptors to monitor.
- **readfds**: specifies the pointer to read file descriptor list.
- **writefds**: specifies the pointer to write file descriptor list.
- **exceptfds**: specifies the pointer to the error file descriptor list.
- **timeout**: deadline for **select()**.

select()

```
void FD_SET (int fd, fd_set *set);  
void FD_CLR (int fd, fd_set *set);  
int FD_ISSET (int fd, fd_set *set); // return: 1 if it's available, else: 0  
void FD_ZERO (fd_set *set);
```

- FD_SET(): Add the file descriptor into the set.
- FD_CLR(): Remove the file descriptor from the set.
- FD_ISSET(): Check if the file descriptor is available.
- FD_ZERO(): Clear the set.

Reference

- [Beej's Guide to Network Programming \(中文\)](#)
- [Beej's Guide to Network Programming \(English\)](#)
- [Linux manual page](#)

Happy coding! ●ω●)ค

TA Email: ntu.cnta@gmail.com